

# Job Scheduling on Midway

Hossein Pourreza


hpourreza@uchicago.edu

April 18, 2017



# Midway

- UChicago's largest super-computer managed by RCC
- A shared resource used by more than 2000 users
- Intended to be used non-interactively



# Job, batch system, and scheduling

- A *job* is an executable code along with input data, parameters, environment variables, output files, and the description of required resources
- Jobs are defined using job scripts
- A batch system (SLURM on Midway) receives job scripts, puts them in a queue, and try to run them according to some rules
- The process of running jobs is called *scheduling*

# Scheduling goals

- Maximize fairness
- Maximize throughput
- Minimize waiting time
- Maximize utilization

# Before submitting a job

- Login to Midway1 or Midway2
  - Depending on where you eventually want to run your job
- Know your default account
  - `rcchelp` user `CNetID`
- Make sure you have enough service units
- Go to the folder from where you want to launch your job
- Write your job script

# Different partitions on Midway

- A partition is a group of compute nodes sharing similar properties
  - Many partitions on Midway
  - Try the `rcchelp sinfo` command for the complete list
- Different partitions have different policies
  - Try the `rcchelp qos` command for the complete list
- Depending on your account type, some of partitions may not be available to run your job

# Public partitions on Midway

<b>Midway1</b>	sandyb
	ivyb
	bigmem
	westmere
	gpu
	mic
	amd
<b>Midway2</b>	broadwl
	broadwl-lc
	bigmem2
	gpu2

# Service Units

- Service Units (SUs) are the measure of computing resources
- In the simplest form, 1 SU is defined as using 1 core for 1 hour
- Each Principal Investigator (PI) requests for SUs to be used by her/his research group
- Use `rcchelp balance` and `rcchelp usage` commands to find out more your balance and usage



# Storage types on Midway

- Home
  - A backed-up, private space usually used to keep your important files
  - Usually not a good place to start your jobs
- Project
  - A backed-up space usually shared between members of a research group
  - Jobs could be started from project unless they require heavy I/O
- Scratch
  - Non backed-up, private space tuned to run large I/O jobs
  - Accessible via the `SCRATCH` environment variable

# Using software on Midway

- Midway uses the module software management tool
- You usually need to load a module before running your code
- The same set of modules should be specified in you job script
- Avoid loading modules in your login startup scripts

# CPU, memory, and walltime

- Serial jobs cannot run faster by asking for more CPUs
  - Ask for as many CPUs as your program can utilize
- Know your job's maximum amount of memory usage
- Know how long it will take for your program to finish
- Asking extra CPUs or memory will unnecessarily burn your SUs without getting your code run faster

# Question 1

- Which of the following statements is correct?
  - My program reads gigabytes of data and writes few megabytes per second. I should start this job from home?
  - I login to Midway1 and then decide which partition to run my job on
  - I always ask 32GB of memory to make sure my job has enough memory to finish
  - I load required modules before run my code on Midway

# Question 2

- True or False?
  - I submit my jobs in midnight to be ahead of other users
  - When there are many jobs on Midway, my program runs slower
  - The more CPU I request, the faster my job will run
  - The Midway scheduler will run my job as soon as possible respecting other jobs' priority

# Before submitting a job

- Login to Midway1 or Midway2
  - Depending on where you eventually want to run your job
- Know your default account
  - rcchelp user CNetID
- Make sure you have enough service units
- Go to the folder from where you want to launch your job
- Write your job script

# A sample job script

```
#!/bin/bash
#SBATCH --job-name=example_sbatch
#SBATCH --output=example_sbatch_%j.out
#SBATCH --error=example_sbatch_%j.err
#SBATCH --account=you_pi_account
#SBATCH --time=00:05:00
#SBATCH --mail-type=BEGIN,END,FAIL
#SBATCH --mail-user=your_email_address
#SBATCH --partition=broadwl
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000

module load python
cd your_working_directory
echo "job started at `date`"

python my_script.py

echo "job finished at `date`"
```

# Submitting a job

- Once you have a job script (saved as `sample_job.sbatch`, for example), you could submit the job as:
  - `sbatch sample_job.sbatch`
- To see if your job has been accepted, run:
  - `squeue -u $USER`



# Multi-threading vs message passing

- Multi-threading is a parallel programming mode utilizing available compute cores of one node
  - Use `--nodes=1` along with `--ntasks=xx` if you are running multi-threaded code
  - A multi-threaded code can NOT use cores from multiple nodes
- Message passing is a parallel programming mode utilizing compute cores from multiple nodes
  - Use `--ntasks=xx` if you running message passing code

# Monitoring jobs

- While your job is running
  - `queue -u $USER`
    - A useful variable in your login startup script
      - `export SQUEUE_FORMAT="% .8i %10P % .8j %10u  
% .15a % .3t %9r % .10l % .10L % .5D % .4C %N"`
  - `scontrol show job job_id`
  - To cancel your submitted job
    - `scancel job_id`
- After your job is finished
  - `sacct -j job_id`
    - A useful variable in your login startup script
      - `export SACCT_FORMAT="jobid,partition,user,  
account%16,alloccpus,node%20,elapsed,totalcp  
u,maxRSS,state"`

# Interactive jobs

- It is sometimes necessary to test your code interactively before submitting your batch job
- To create an interactive session, use the `sinteractive` command:
  - `sinteractive -p sandyb --ntasks=1 --mem-per-cpu=2000 --time=00:30:00`
- Once you get a session, you could load modules (if applicable) and run your code interactively

# GPU-enabled jobs

- Midway has some special purpose nodes dedicated to run jobs using Graphical Processing Units (GPUs)
- Your program should be GPU enabled to take advantage of these nodes
- To request a GPU node with one GPU card, use:
  - `#SBATCH --partition=gpu`
  - `#SBATCH --gres=gpu:1`
  - The `--partition=gpu` and `--gres=gpu:1` tags can also be used with the `sinteractive` command

# Job arrays

- The job array in SLURM is an efficient way to submit many similar jobs

```
#!/bin/bash
#SBATCH --job-name=arrayJob
#SBATCH --output=arrayJob_%A_%a.out
#SBATCH --error=arrayJob_%A_%a.err
#SBATCH --array=1-6
#SBATCH --time=01:00:00
#SBATCH --partition=sandyb
#SBATCH --ntasks=1
#SBATCH --mem-per-cpu=4000

module load python
# Print this sub-job's task ID
echo "My SLURM_ARRAY_TASK_ID: " $SLURM_ARRAY_TASK_ID

# Do some work based on the SLURM_ARRAY_TASK_ID
python my_script.py $SLURM_ARRAY_TASK_ID
```

# Reducing the number of serial jobs

```
#!/bin/sh
#SBATCH --time=00:05:00
#SBATCH --ntasks=4

module load python
module load parallel

srun="srun --exclusive -N1 -n1"

parallel="parallel --delay .2 -j $SLURM_NTASKS --joblog
runtask.log --resume"

$parallel "$srun python my_script.py {1} > runtask.{1}" :::
{1..16}
```

# Conclusion

- Midway is a world-class super computer accessible to all UChicago researchers
- To be able to run your code on Midway, you need to submit a job
- Selecting accurate parameters in your job submission script can reduce your wait time and SU usage

# Useful links

- RCC official website
  - <https://rcc.uchicago.edu>
- Midway user documentation
  - <https://rcc.uchicago.edu/docs/using-midway/index.html>
- Official SLURM commands documentation:
  - <https://slurm.schedmd.com/sbatch.html>
  - <https://slurm.schedmd.com/scontrol.html>
  - <https://slurm.schedmd.com/squeue.html>
  - <https://slurm.schedmd.com/sacct.html>